

12-1-15
P.1

2

OPERATING SYSTEMS

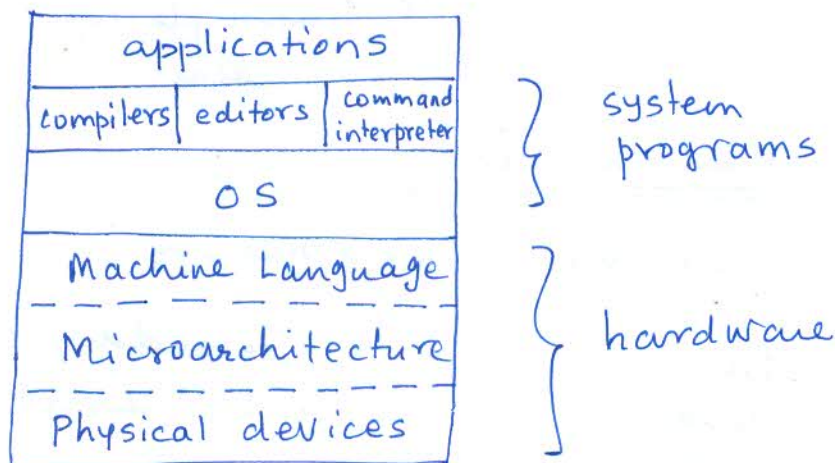
Mondays 12:00 noon
+
Thursdays to
1:25 pm

Text book -

Modern Operating Systems 2nd ed.

by A. S. Tanenbaum
PKI

INTRODUCTION



1.1.1 The Operating System as an
Extended Machine

1.1.2 as a
Resource Manager

brief history -

technology → mainframes / minis

CPU speed vs. I/O speed

multiprogramming

timesharing

UNIX - POSIX - LINUX

Windows / GUI

Network OS vs Distributed OS

↘ problems with
multiple copies of
data

Hardware review

CPU ↔ cache ↔ memory

1
fetch
decode
execute

{ secondary storage
I/O devices

{ DMA
programmed
I/O

interrupt mechanism

interrupt handling

buses

controllers



bootstrapping

8-1-15
p.1

interrupts - why?

3

* Fig 1-10
(slight change
in (b))



DMA controller may
come into play


15-1-15
1.5

OPERATING SYSTEM CONCEPTS

process program in execution
address space of process ["core image"]
process table
creation / termination
parent process - child process

(p.2)

process tree Fig 1-12




⇒ need two flavours
of process creation
"blocking"
"non-blocking"

* interprocess communication

user ID vs. process ID

Deadlocks

- to be distinguished from simply
"running out of resources"



Memory management -

Secondary storage

SIMULATION

File system

19-1-15
(P.1)

1.6 SYSTEM CALLS

- * interface between user programs & OS
- * difference between usual function calls (a.k.a. procedure calls) and system calls → the latter "trap" to OS kernel

e.g. `count = read (fd, buffer, nbytes)`

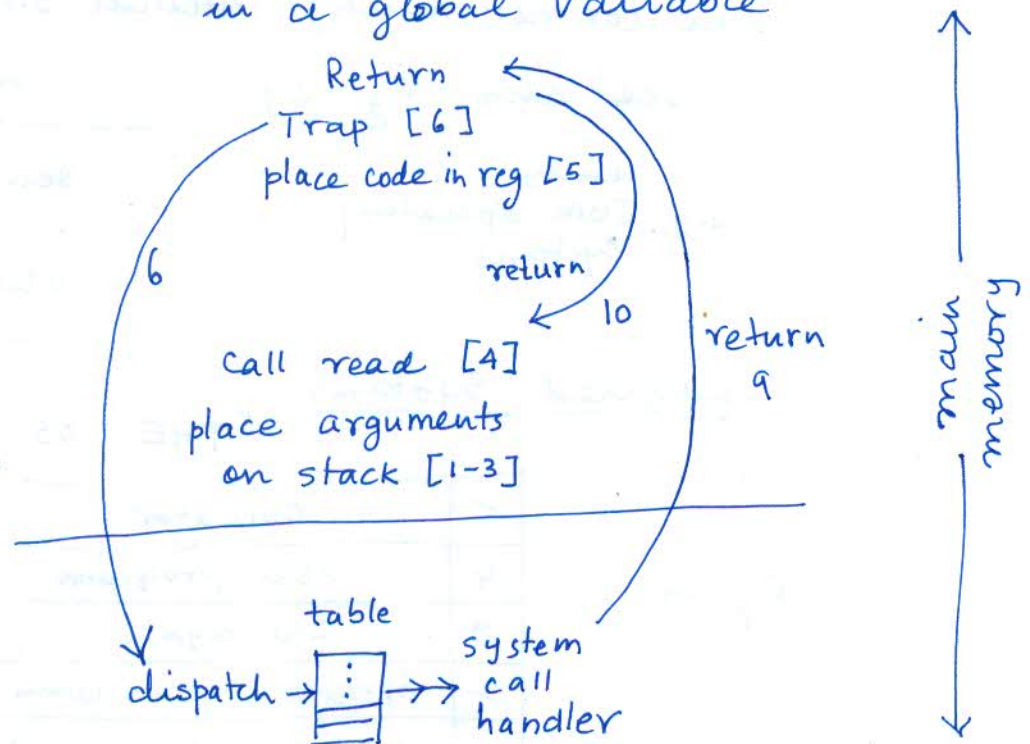
if EOF, then `count < nbytes`

if error, `count = -1` and error code in a global variable

Fig 1-17

user space

kernel



(p.2)

* POSIX - IEEE standard

Fig. 1-18 - Self study

Fig. 1-23 - Win32 API (subset)

Recall - OS as "Extended Machine"

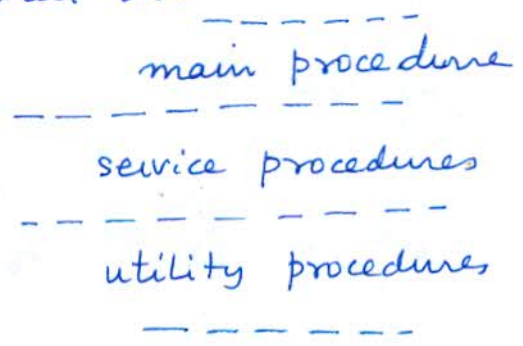
1.7 Operating System Structure

Monolithic Systems

- collection of procedures
- one big object file
- no (or not much) internal structure

see also Fig 1-24

e.g. ^{early} IBM operating systems



Layered Systems


"THE" OS - E.W. Dijkstra
1968

Fig 1-25

5	Operator
4	User programs
3	I/O mgmt
2	operator-process comm.
1	memory + drum mgmt
0	processor allocation

← multi-programming

Virtual machines



← "identical copies"
of hardware
[trap to h/w
as needed]

Fig 1-26 shows VM/370 with CMS

extension of this concept

Exokernels

each defined VM has specified
subset of hardware resources

[not a copy of the "whole m/k"]

Client-Server model

- on one system [with "microkernel"]
- on a distributed system

* recall difference
between Network OS
and Distributed OS

22-1-15
(P.1)

Chap.2 PROCESSES

"most central concept"

- resources must be fully utilized
- multiple activities in progress at any one time
- "pseudo parallelism" as opposed to (h/w) multiprocessing
- process - conceptual model (sequential) which helps deal with parallelism






Fig 2-1
(a)




four programs
in memory

(b)



four independent,
sequential processes
each with "own PC"

(c)



(p.2)

- CPU switches back and forth between processes
- rate of progress of any process depends on total system load
- but RESULTS DO NOT!

* [In certain special systems, real-time constraints are important.]

need Real-Time Operating System

Q: What defines process "state"?

Process creation

- At:
- system initialization
 - "create process" system call (executed by another process)
 - user request
 - initiation of a "batch job"
- yes, but !!!

BTW - "batch" processing
vs
"continuous" processing

background vs. foreground processes

\ "daemons"

22-1-15

P.3

20

UNIX - "fork" system call

- creates copy (clone) of calling process

[usually followed by "exec"]

- each process has its own address space

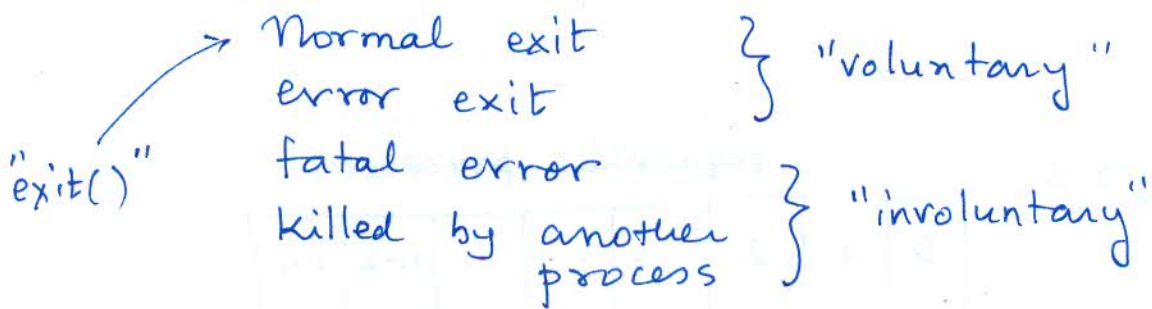
[code can be shared] * → "pure" or "reentrant" code

note

resources

needed for process creation must be available for the system call to succeed.

Process termination



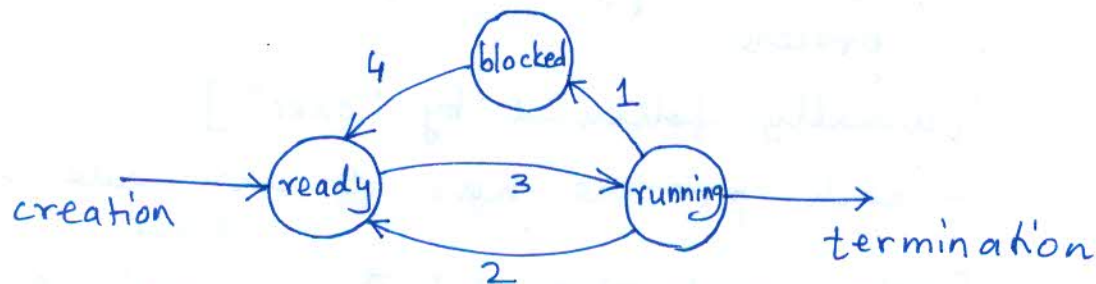
Process hierarchy

- easily seen in UNIX / LINUX ...
- not in WINDOWS

i.e. not an indispensable feature of an OS, but a neat and convenient one


↑
can prove useful in managing processes

Process states



- 1: process blocks for I/O or another event
- 2: scheduler picks another process
- 3: " " this "
- 4: I/O completes or event occurs

Fig 2-3



← lowest-layer of OS handles interrupts & scheduling

Such a system would be described as a "process-structured" operating system.

29-1-15

(P.1)

22

Implementation of Processes

Process Table - data structure e.g. array of structures ["process control blocks"]

Fig 2-4 Typical Process Table entry

Process Mgmt - registers, PC, PSW, SP ...
process state, priority
scheduling parameters
ID, parent ID, process group
signals
Time started, CPU time used

Memory Mgmt - Pointers to text, data and stack segments
[UNIX terminology]*

File Mgmt - Root directory
Working "
File descriptors
User ID, Group ID

note: various fields listed above
help with* resource mgmt
or required for

Fig 2-5 (simplified) what an OS does upon an interrupt

1. Hardware stacks PC
2. " loads new PC from interrupt vector
3. A.L. procedure saves registers
4. " " " sets up new stack
5. ISR runs [can be in C]
6. Scheduler decides which process is to run next
- 7-8 A.L. procedure starts up (resumes) selected process

interrupted process and selected process (step 6) may be different

* hardware dependent code is written in Assembly language [small percentage of OS code]

29-1-15
(P.3)

2.3

THREADS

process provides
(so far)

(a)

resource grouping

single thread of
execution

(b)

often it is useful - from app_n
point of view - to provide (a)
with multiple threads (of execution)

for execution, each thread must
have its own PC, registers, stack, state
& SP

threads share all other resources

- address space, open files,
accounting, access rights.....

- sometimes known as "light-weight processes"
- capability known as "multithreading"

* self study - Fig 2-6, Fig 2-7

- at the time of creation, a process
(usually) starts off as a single
thread

(p.4) as needed, "system calls" such
 as thread-create
 thread-exit are invoked
 thread-wait
 thread-yield

"careful thought and design
 are needed to make multithreaded
 programs work correctly"

Fig 2-9 a word-processor with 3 threads

- user interaction
- reformatting (b/g)
- periodic saving
 of file to disk (b/g)

Figs 2-10, 2-11

dispatcher thread:

```
while (TRUE) {
    get-next-request(&buf);
    handoff-work(&buf);
}
```

9-2-15

(P.1)

2.3 INTERPROCESS COMMUNICATION

25

need { explicit
implicit ← sharing of system resources

three types of communication
coordination]

- (i) one process passes information to another
- (ii) processes should not "get into each other's way" during critical activities
- (iii) logical dependencies

Issues are applicable to threads as well, except that:

- (a) they have common address space
- (b) they are part of a common* application
* [multithreaded]
(and presumably correctly programmed as such)

P.2

2.3.1 Race conditions

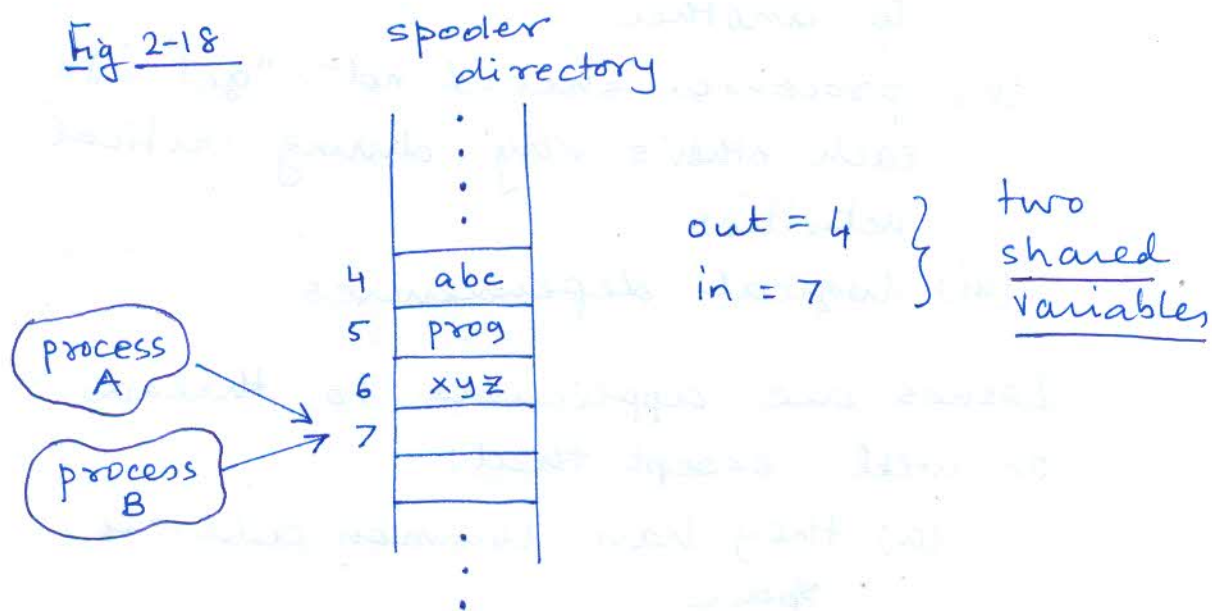
→ when two or more processes access shared storage (main memory or on a file)

Example -

Print spooler

spooler directory - printer daemon

Fig 2-18



two or more processes are reading and writing shared data, but the final result* depends on which process runs when

[recall that interrupts and the scheduler are running independently]

correctness
of the final
result

9.2-15

p.3

2.3.2

Critical regions

25

→ two (or more) processes should not read & write shared data "at the same time" (i.e. in any possible sequence of reads & writes)

→ i.e. if one process is changing the value of a shared variable, other processes should be excluded from doing the same thing until the first one has "completed its work on the shared variable"

→ concept of critical region* helps us to formulate a solution

[or critical section]

→ no two processes should be in

① their respective critical region at the same time

↑ Also - ② no assumptions about speeds or number of processors

"good solution" → ③ no process outside its CR may block other processes

→ ④ no process should have to "wait forever" to enter its CR

p. 4




Fig 2-19

2.3.3

Mutual exclusion with "busy waiting"

- * disabling interrupts
- * lock variables
- * strict alternation — violates condition ③ seen above
- * Petersen's solution

"busy waiting" is not really acceptable in a good solution

\rightarrow must look for a better solution

12-2-15
P.1

TSL machine instruction

20

TSL \rightarrow stands for "Test & set Lock"

TSL RX, LOCK

indivisible action: $\begin{cases} RX \leftarrow (LOCK) \\ (LOCK) \leftarrow 1 \end{cases}$

* LOCK is
a memory
location

\downarrow
even on a multiprocessor system

Fig 2-22

enter_region: TSL Ri, LOCK
CMP Ri, #0
JNE enter_region
RET

leave_region: MOVE LOCK, #0
RET

each process contending for a shared resource must:

- (a) call enter-region at point of entry to ^{its} CR
- (b) call leave-region at point of exit from ^{its} CR

But note: busy waiting has not yet been removed from the solution!

P.2

The "Producer-Consumer problem"

- two processes share a fixed size buffer (which can store N items)

Consider Fig 2-23

#define N 100

int count = 0;

void producer(void)

{

int item;

while (TRUE) {

item = produce();

if (count == N) sleep();

insert-item(item);

count++;

if (count == 1) wakeup(consumer);

}

{

void consumer(void)

:

if (count == 0) sleep();

item = remove-item();

count--;

if (count == N-1) wakeup(producer)

consume-item(item)

:

(p.3) But there is still a race condition -

consumer reads 'count' when it is 0

* — process switch — *

producer inserts item, sets count to 1
sends "wakeup" to consumer

* — process switch — *

consumer goes to sleep
(as per code)

Problem - the wakeup call sent by
the producer is "lost"!

2.3.5 Semaphores E.W. Dijkstra (1965)

→ solve the problem of "lost wakeups"

→ integer valued variables with two
operations defined up()
down()

→ no busy waiting

→ if "no stored wakeups", process is
blocked [as part of down()]

→ down() & up() \Rightarrow ATOMIC

P.4

Solution of Producer-Consumer problem
using semaphores -

```
#define N 100
typedef int semaphore;
semaphore mutex = 1;
semaphore empty = N;
semaphore full = 0;
```

producer
loop

```
item = produce();
down(&empty);
down(&mutex);
insert_item(item);
up(&mutex);
up(&full);
```

consumer
loop

```
down(&full);
down(&mutex);
item = remove_item();
up(&mutex);
up(&empty);
consume(item);
```